

Notes on Computability & Complexity

XIN CHEN chenxin_hello@outlook.com $Q_{quality} = \int(K, P, t)$

latest update: December 22, 2024

May the force of P and NP be with you.

Contents

1	Basic concepts	3
1.1	logarithms	3
1.2	Strings	3
1.3	Representations	4
1.4	Big-Oh notation	4
2	Computational models	6
2.1	Turing machines	6
2.2	Running Time	7
2.3	Machines as Strings and the Universal Turing Machine	7
2.4	Variants of Turing machines	8
2.4.1	Turing machines with alphabet $\{0, 1, \sqcup, \triangleright\}$	8
2.4.2	Single tape Turing machines	8
2.4.3	Bidirectional single tape Turing machines	8
2.5	Other computational models*	8
2.5.1	URM: the unlimited register machine	8
3	Uncomputability	9
4	Complexity Classes	10
4.1	Hardness and Completeness	10
5	The Class P	10
6	The Class NP, $coNP$ and DP	11

Citation testing: [1]

One of the important scientific advances in the first half of the twentieth century was that the notion of “computation” received a much more precise definition.

At roughly 50 years (1970s – 2020s), complexity theory is still an infant science, and many important results are less than 30 years old.

1 Basic concepts

1.1 logarithms

Definition 1.1 (Logarithms) Let a, b be two positive real numbers, and $a \neq 1$. The *logarithm of b to the base a* , denoted by $\log_a b$, is the unique real number x such that $a^x = b$. That is,

$$\log_a b = x \quad \text{iff} \quad a^x = b.$$

⊣

The logarithm function is the inverse of the exponential function. The logarithm function is defined for $a > 0$, $a \neq 1$, and $b > 0$. The most commonly used bases are $a = 2$, $a = e$, and $a = 10$. The base $a = 2$ is used in computer science, $a = e$ is used in calculus, and $a = 10$ is used in engineering and science. The base $a = 10$ is called the *common logarithm*, and is denoted by $\lg b$.

Table 1: Commonly used bases for logarithms

Base	Notation
2	$\log b$
e	$\ln b$
10	$\lg b$

$$\log_a b \rightsquigarrow a^\square = b$$
$$e \approx 2.718$$

Table 2: Identities of logarithms

Identity	Formula
Product	$\log_b(xy) = \log_b x + \log_b y$
Quotient	$\log_b(\frac{x}{y}) = \log_b x - \log_b y$
Power	$\log_b x^r = r \cdot \log_b x$
Root	$\log_b \sqrt[r]{x} = \frac{1}{r} \cdot \log_b x$

1.2 Strings

If S is a *finite* set, called *alphabet set*, then a *string* over S is a finite ordered tuple of elements from S .

We will typically consider the *binary* alphabet $2 = \{0, 1\}$.

$$S^0 = \{\epsilon\}$$

$S^* = \bigcup_{n \geq 0} S^n$ is the set of all strings over S .




Figure 1: Plots of logarithm functions, with three commonly used bases, from wikipedia

The *concatenation* of strings x, y is denoted by $x \frown y$, $x \circ y$, or simply xy .

x^k denotes the concatenation of k copies of x for $k \geq 1$. For example, 1^3 is ‘111’.

The length of a string x is denoted by $|x|$.

1.3 Representations

we implicitly identify any function f whose domain and range are not strings with the function

$$g: \{0, 1\}^* \rightarrow \{0, 1\}^*$$

that given a representation of an object x as input, outputs the representation of $f(x)$.

1.4 Big–Oh notation

Definition 1.2 (Big–Oh notation) If f, g are two functions over \mathbb{N} , then we say that

- (1) $f = O(g)$ if there exist a constant c such that $f(n) \leq c \cdot g(n)$ for every sufficient large n .
- (2) $f = \Omega(g)$ if $g = O(f)$.
- (3) $f = \Theta(g)$ if $f = O(g)$ and $g = O(f)$.
- (4) $f = o(g)$ if for $\forall \kappa > 0$, $f(n) \leq \kappa \cdot g(n)$ for every sufficient large n .
- (5) $f = \omega(g)$ if $g = o(f)$.

To emphasize the input parameter, we often write $f(n) = O(g(n))$ instead of $f = O(g)$, and use similar notation for $o, \Theta, \Omega, \omega$. ⊣

Example 1.3 Here are some examples of big–Oh notation:

- (1) If $f(n) = 100n \log n$ and $g(n) = n^2$, then $f = O(g)$.

(2) If $f(n) = 100n^2 + 24n + 2\log n$ and $g(n) = n^2$, then $f = O(g)$ and $g = O(f)$.

+

2 Computational models

This section introduces some of the most important computational models in the history of computer science and explains *why it doesn't matter*.

Uncomputability has an intimate connection to Gödel's famous Incompleteness Theorem.

2.1 Turing machines

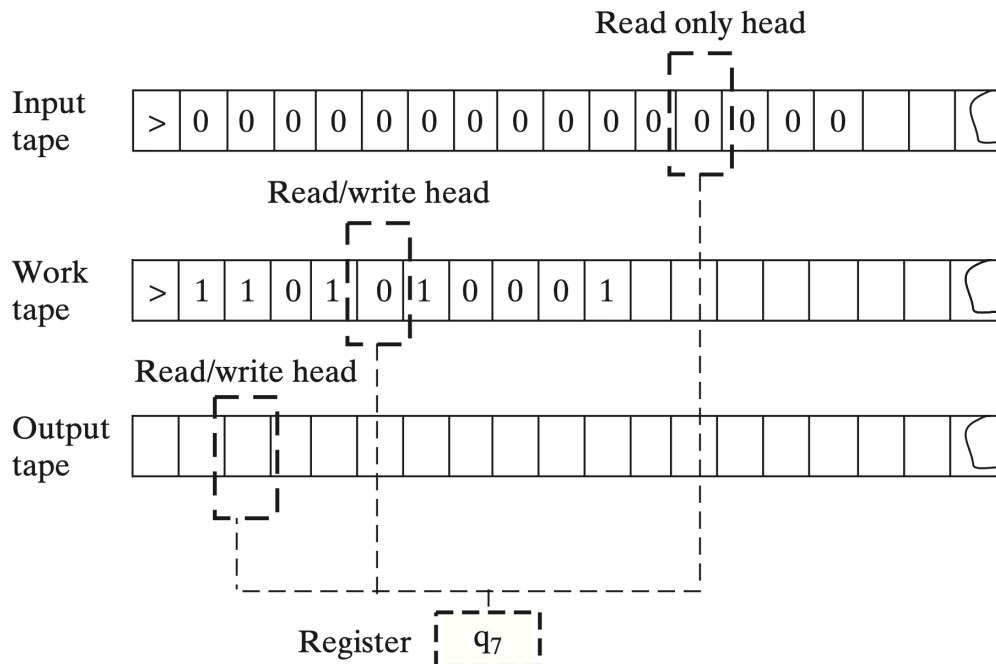


Figure 2: A snapshot of the execution of a three-tape Turing machine M with an input tape, a work tape, and an output tape, see [1, p. 11]

A *tape* is an infinite one-directional line of cells, each of which can store a symbol from a finite set called *alphabet*. Each tape is equipped with a *tape head* that can potentially read or write symbols to the tape one cell at a time.

alphabet $\{0, 1\} \cup \{\sqcup, \triangleright, \}$

state set Q , contains two distinguished states: the start state q_{start} and the halting state q_{halt}

Example 2.1 (palindrome [回文]) A *palindrome* is a string that reads the same forwards and backwards. The language of palindromes over the binary alphabet is

$$PAL = \{w \in \{0, 1\}^* \mid w = w^R\}$$

where w^R denotes the reverse of w . For example, 00100 is a palindrome, and clearly $\{\epsilon, 0, 1\} \subseteq PAL$.

A Turing machine that computes *PAL* within less than $3|x|$ steps for any input x .

Our TM M will use three tapes (the input, work and output tape) and the alphabet $\{\sqcup, \triangleright, 0, 1\}$:

- 1: Input x (where $|x| = n$)
- 2: Copy x to the work tape $(n \text{ steps})$
- 3: Move the input-tape head the beginning of x $(n \text{ steps})$
- 4: Move the input-tape head to the right while moving the work-tape head to the left.
- 5: **if** at any moment the machine observes two different symbols **then**
- 6: **reject** and output 0
- 7: **else**
- 8: **accept** and output 1
- 9: **end if** $(\leq n \text{ steps})$

–

2.2 Running Time

Any nontrivial computational task requires at least reading the entire input, we count the number of basic steps as a function of the input length.

Definition 2.2 (Computing a function and running time) Given $f: \{0, 1\}^* \rightarrow \{0, 1\}^*$, $T: \mathbb{N} \rightarrow \mathbb{N}$ and a Turing machine M , we say that

- (1) M *computes* f if for every $x \in \{0, 1\}^*$, whenever M is initialized to the start configuration on input x , then it halts with $f(x)$ on the output tape.
- (2) M *computes* f *in* $T(n)$ -*time* if its computation on every input x requires at most $T(|x|)$ steps.

–

2.3 Machines as Strings and the Universal Turing Machine

It is almost obvious that *we can represent a Turing machine as a string*: Just write the description of the TM on paper, and encode this description as a sequence of zeros and ones. This string can be given as input to another Turing machine.

2.4 Variants of Turing machines

2.4.1 Turing machines with alphabet $\{0, 1, \sqcup, \triangleright\}$

2.4.2 Single tape Turing machines

2.4.3 Bidirectional single tape Turing machines

2.5 Other computational models*

2.5.1 URM: the unlimited register machine

3 Uncomputability

There exist functions that cannot be computed within any finite amount of time.

4 Complexity Classes

4.1 Hardness and Completeness

For any complexity class \mathcal{C} , a decision problem D is said to be \mathcal{C} -hard if

$$\forall F \in \mathcal{C}, F \leq_p D.$$

If, in addition $D \in \mathcal{C}$ then D is said to be \mathcal{C} -complete.

So the class of NP-complete problems are those problems in NP to which any other problem in NP can be polynomially reduced.

- (1) deciding if a propositional formula has a model (SAT).
- (2) deciding if a graph has a path that contains every vertex exactly once (a variant of the so-called Travelling Salesperson Problem).
- (3) deciding if a given argument is acceptable w.r.t. Dung's stable semantics.

5 The Class P

A *complexity class* is a set of functions that can be computed within given resource bounds.

Definition 5.1 (The class DTIME) The class $\text{DTIME}(f(n))$ is the set of all functions that can be computed by a deterministic Turing machine in $O(f(n))$ steps. \dashv

The D in the notation DTIME refers to “deterministic”.

Definition 5.2 (The class P) The class P is the union of all complexity classes $\text{DTIME}(n^c)$ for all c . In other words,

$$P = \bigcup_{c \geq 1} \text{DTIME}(n^c).$$

\dashv

P: polynomial.

Example 5.3 (Graph connectivity) Graph connectivity problem:

input: a graph G and two vertices s, t ,

inquiry: whether there is a path from s to t .

Using *depth-first search* (DFS), we can solve this problem in $O(|V| + |E|)$ steps. \dashv

.....

A problem is viewed as having an efficient algorithmic solution if it can be placed into the class **P** (*polynomial-time*) of all problems that have a polynomial-time algorithm, i.e., an algorithm that for each instance x (of size $|x|$) produces its answer after at most $|x|^k$ steps, for a fixed constant k .

P is a rather coarse-grained class, an important subclass we will consider is **L** (*logarithmic space*), which consists of the problems that can be solved in logarithmic space (not counting input and output) and polynomial-time.

We consider problems in the classes **L** and **P** to be computationally *tractable*.

6 The Class **NP**, **coNP** and **DP**

.....
 Quantified boolean formula (QBF)

$$Q_1x_1Q_2x_2\cdots Q_kx_k\varphi(x_1,\dots,x_k)$$

where $Q_i \in \{\forall, \exists\}$ and \exists is followed by \forall

References

- [1] Sanjeev Arora and Boaz Barak. *Computational Complexity: A Modern Approach*. Cambridge: Cambridge University Press, 2009. xxiv+579. ISBN: 978-0-511-53381-5. URL: <https://doi.org/10.1017/CB09780511804090> (cit. on pp. 2, 6).